

Certified Kubernetes Application Developer (CKAD)

By Damian Igbe, PhD

Day 4

Encrypting with SOPS and Age

With the ever-growing need to safeguard secrets and sensitive data, the challenge is to find tools that offer a balance between security and usability. One such pairing that stands out is [Mozilla's SopS](#) in combination

with [AGE for encryption](#). In this blog post, we'll dive deep into the benefits of this duo, touching on simplicity, version control, encryption, and even offer a hands-on example.

What is SopS?

SopS is a secret management tool developed by Mozilla. It's designed for encrypting secrets in files that can be stored in a version control system like GIT. Rather than encrypting the entire file, SopS encrypts only the specific parts of a file that contain secrets, leaving the rest of the file (like structure and comments) in plain text. This means that you can still track changes and maintain version history without compromising security. Furthermore, SopS supports a variety of secret backends i.e. cloud KMS and Hashicorp vault. SopS additionally works in a binary mode where full files can be encrypted.

Introducing AGE

AGE (Actually Good Encryption) is a simple, modern and secure encryption tool with small explicit threat models. Built on the foundations of strong cryptographic primitives, AGE focuses on being simple to use while retaining powerful encryption capabilities.

Why Pair SopS with AGE?

1. **Simplicity:** One of the primary benefits of using AGE is its simplicity. With a straightforward CLI and easy-to-understand encryption and decryption commands, it reduces the chance of human error. When integrated with SopS, you get the combined power of both tools intuitively.
2. **Version Control Compatibility:** Since SopS encrypts only the secrets and not the entire file, it plays well with version control systems. You can visualize diffs, track changes, and maintain an audit trail of modifications without revealing the actual secrets.
3. **Strong Encryption:** AGE boasts robust encryption capabilities. It uses modern encryption techniques, ensuring that your data remains confidential and safe from potential breaches.

Code Example: Encryption and Decryption with SopS and AGE

First, generate an AGE keypair:

```
age-keygen -o age.key
```

This will produce a public key string that you can use for encryption and a private key saved in age.key for decryption.

Next, use SopS to encrypt a file using the AGE public key:

```
sops --age age_public_key -e secrets.yaml > secrets.enc.yaml
```

Where age_public_key is the public key string generated in the first step and secrets.yaml is your plaintext file.

To decrypt the file, use:

```
sops --age age.key -d secrets.enc.yaml
```

The encrypted file (secrets.enc.yaml) can be safely stored in version control, while the original secrets remain safe and sound.

Lab Walkthrough

Objective

- Use **SOPS** (secrets OperationS) to manage secrets.
- Encrypt secrets using **AGE (Actually Good Encryption)** encryption.
- Store secrets securely in **Kubernetes** and access them via **Kubernetes Secrets**.
- Decrypt secrets dynamically in your Kubernetes deployments.

Prerequisites

1. **Kubernetes** cluster
 2. **kubectl** command-line tool installed and configured to interact with your cluster.
 3. **SOPS** installed: A tool to manage encrypted secrets files.
 4. **AGE** installed: A simple and secure encryption tool.
 5. **Helm** installed (optional, for better management).
 6. Basic knowledge of Kubernetes.
-

Step 1: Install sops and age

Install sops

Linux:

```
curl -Lo sops https://github.com/mozilla/sops/releases/download/v3.7.3/sops-v3.7.3-linux-amd64
chmod +x sops
sudo mv sops /usr/local/bin/
```

Install age

Linux:

```
curl -Lo age.tar.gz https://github.com/FiloSottile/age/releases/download/v1.0.0/age-linux-amd64.tar.gz
tar -xvf age.tar.gz
chmod +x age
sudo mv age /usr/local/bin/
```

Step 2: Generate an Age Key Pair

AGE encrypts secrets using a public/private key pair. Let's generate the keys.

```
age-keygen -o key.txt
```

This command generates a new key pair and outputs it to key.txt. The public key will be used for encryption, and the private key will be used for decryption.

You will see output like:

```
# created: 2025-03-12T00:00:00Z
# public key: age1xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
<private_key_content>
```

Store the private key securely, as it will be required later for decryption. You can move it to ~/.sops/

Export the key to an ENV and ensure that it works

```
SOPS_AGE_KEY_FILE=~/.sops/key.txt
echo $SOPS_AGE_KEY_FILE
~/.sops/key.txt
```

you can reference here for different environments <https://github.com/getsops/sops>

Step 3: Encrypt a Secret Using AGE

Now that you have your public key (from the key.txt file), let's encrypt a secret. For example, let's say you want to encrypt a Kubernetes password.

1. Create a file with the secret, e.g., db-password.txt:

```
echo "mysecretpassword" > db-password.txt
```

2. Encrypt the file using AGE:

```
age -r <public_key> -o db-password.txt.age db-password.txt
```

Where <public_key> is the public key from the key.txt file generated earlier. After running this, the file db-password.txt.age will be the encrypted version of your password.

You can delete the plaintext secret file for security:

```
rm db-password.txt
```

Step 4: Create a Kubernetes Secret for the Encrypted Data

We can now create a Kubernetes secret containing the encrypted data. This ensures the secret is stored securely in Kubernetes.

```
kubectl create secret generic db-password --from-file=db-password.txt.age
```

This command creates a Kubernetes secret named db-password, which contains the encrypted password file (db-password.txt.age).

Verify the secret is created:

```
kubectl get secret db-password -o yaml
```

You should see something like:

```
apiVersion: v1
data:
  db-password.txt.age: <encrypted_content_here>
kind: Secret
metadata:
  name: db-password
  namespace: default
type: Opaque
```

The db-password.txt.age will contain the encrypted password.

Step 5: Use SOPS for YAML-based Secret Management

Instead of manually encrypting secrets, you can use SOPS to manage your secrets as YAML files that are both human-readable and encrypted. This is very useful when working with Kubernetes manifests.

1. Create a secret manifest, e.g., secret.yaml:

```
apiVersion: v1
kind: Secret
metadata:
  name: db-password
type: Opaque
data:
  db-password.txt.age: <Base64 encoded encrypted secret>
```

2. Encrypt the YAML file using SOPS:

```
sops --encrypt --age <public_key> secret.yaml > secret-encrypted.yaml
```

The secret-encrypted.yaml file will now contain the encrypted version of your secrets.

To decrypt it:

```
sops --decrypt secret-encrypted.yaml
```

Step 6: Decrypt Secrets at Runtime with a Kubernetes Deployment

In most cases, you will not store plaintext secrets directly in your code. Instead, you can use Kubernetes to decrypt them when needed.

1. Create a Kubernetes Deployment manifest that will access the secret:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: secret-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: secret-app
  template:
    metadata:
      labels:
        app: secret-app
    spec:
      containers:
        - name: secret-app
          image: nginx
          env:
            - name: DB_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: db-password
                  key: db-password.txt.age
```

This Deployment will mount the db-password.txt.age file from the Kubernetes secret as an environment variable.

To apply the deployment:

```
kubectl apply -f deployment.yaml
```

Step 7: Use SOPS and AGE in CI/CD Pipelines

You can automate the process of encrypting and decrypting secrets using CI/CD tools like **Jenkins**, **GitLab CI**, or **GitHub Actions**. Here's how you can implement the decryption step dynamically during your pipeline execution.

1. **Decrypt secrets in your CI/CD pipeline:**
 - Store the private key securely (preferably using your CI/CD secrets manager).
 - Use age to decrypt the secrets and inject them into your application as environment variables.

```
age -d -i /path/to/private/key.db-password.txt.age
```

2. Make sure you handle the private key and the decrypted secrets securely to avoid accidental exposure.
-

Step 8: Clean Up

Finally, clean up your Kubernetes resources when you are done:

```
kubectl delete secret db-password  
kubectl delete deployment secret-app
```

Conclusion

This lab demonstrates how to securely manage and encrypt Kubernetes secrets using **SOPS** and **AGE**. It walks you through the process of encrypting secrets, storing them in Kubernetes, and making them available to your applications securely. This workflow ensures your sensitive data is protected at rest, while still being accessible to your applications when needed.

Domain 4.8: Understand Application Security (SecurityContexts, Capabilities, etc.)

Application security in Kubernetes focuses on securing containers and their environments to minimize vulnerabilities and risks. This involves configuring and controlling various security features to ensure that applications run with the least privilege, restrict unnecessary permissions, and protect the integrity of both the host and the containers.

In this section, we'll cover:

- **SecurityContext:** What it is, and how to use it for controlling security settings at the Pod or container level.
- **Linux Capabilities:** Understanding and configuring Linux capabilities for containers to limit the permissions granted to a container.
- **Pod Security Policies (PSP):** How to enforce security policies at the Pod level (if enabled).
- **Other Key Security Features:** read-only file systems, seccomp, and AppArmor.

1. What is a SecurityContext?

A **SecurityContext** is a Kubernetes resource used to define security-related settings for Pods and containers. You can specify a **SecurityContext** at two levels:

1. **Pod-level SecurityContext:** This applies to all containers in the Pod.
2. **Container-level SecurityContext:** This applies to individual containers in the Pod. **SecurityContext** allows you to configure various security settings such as user and group IDs, privilege escalation, and more. By configuring the **SecurityContext**, you can ensure that containers run with the appropriate security configuration, reducing the risk of exploitation.

Common SecurityContext Settings

- **runAsUser:** Specifies the user ID (UID) to run the container as.
- **runAsGroup:** Specifies the group ID (GID) to run the container as.
- **fsGroup:** The group ID to apply to files created by the container in volumes.
- **privileged:** If true, gives the container extended privileges. By default, containers do not run in privileged mode.
- **allowPrivilegeEscalation:** If set to false, it prevents a container from gaining more privileges than its parent process.
- **readOnlyRootFilesystem:** If set to true, it mounts the root filesystem as read-only.
- **capabilities:** Controls which Linux capabilities to add or drop from the container.

Example of Pod-level SecurityContext

apiVersion: v1

kind: Pod

metadata:

name: secure-pod

spec:

securityContext:

runAsUser: 1000

runAsGroup: 3000

fsGroup: 2000

containers:

- name: secure-container

image: nginx

securityContext:

allowPrivilegeEscalation: false

readOnlyRootFilesystem: true

In this example:

- The Pod-level SecurityContext specifies that the Pod will run with a user ID of 1000, group ID of 3000, and any files created in volumes will be associated with the group ID 2000.
- The Container-level SecurityContext ensures the container cannot escalate privileges (allowPrivilegeEscalation: false) and the root filesystem is mounted as read-only (readOnlyRootFilesystem: true).

2. Linux Capabilities

Linux capabilities are a set of fine-grained access controls that allow you to grant or remove specific privileges from processes. Containers by default have a limited set of Linux capabilities to minimize security risks. However, some applications require additional capabilities. Key Capabilities

- NET_ADMIN: Allows the process to configure networking interfaces.
- SYS_ADMIN: Grants a range of system-level administrative capabilities, such as mounting file systems.
- DAC_OVERRIDE: Allows overriding file read, write, and execute permissions.

How to Modify Capabilities

In Kubernetes, you can add or remove Linux capabilities from a container using the securityContext.capabilities field. For example, if you want to add the NET_ADMIN capability, you can do so like this:

apiVersion: v1

kind: Pod

metadata:

name: capability-pod

```
spec:
  containers:
  - name: capability-container
    image: nginx
    securityContext:
      capabilities:
        add:
        - NET_ADMIN
```

In this example, the container is granted the NET_ADMIN capability, which allows it to modify networking settings.

Best Practice:

Always try to remove unnecessary capabilities to adhere to the principle of least privilege. Adding capabilities should only be done when absolutely necessary.

3. Pod Security Policies (PSP)

Pod Security Policies (PSPs) are a set of policies that control the security features available to Pods in a cluster. PSPs allow you to define rules around:

- What security context settings are allowed.
- What privileges containers can request.
- Restrictions on the use of host network, volumes, and namespaces.

Key PSP Settings:

- privileged: Whether privileged mode containers are allowed.
- runAsUser: Controls which user IDs are allowed to run.
- allowPrivilegeEscalation: Defines if privilege escalation is allowed.
- hostNetwork: Restricts Pods from using the host network.
- hostPID: Restricts Pods from using the host's process ID namespace.
- readOnlyRootFilesystem: Enforces read-only root filesystem for Pods.

PSP Example:

```
apiVersion: policy/v1beta1
```

```
kind: PodSecurityPolicy
```

```
metadata:
```

```
  name: restricted-psp
```

```
spec:
```

privileged: false

runAsUser:

rule: MustRunAsNonRoot

allowPrivilegeEscalation: false

volumes:

- configMap

- secret

hostNetwork: false

readOnlyRootFilesystem: true

In this example, the PSP:

- Disallows privileged containers.
- Requires containers to run as non-root.
- Disallows privilege escalation.
- Restricts volumes to only ConfigMaps and Secrets.
- Disallows the use of the host network.
- Enforces read-only root filesystem for containers.

Note: PSP is deprecated and will be removed in Kubernetes 1.25+. It is being replaced with Pod Security Admission (PSA).

4. Other Security Features

a) Seccomp Seccomp (Secure Computing Mode) is a Linux kernel feature that allows you to filter system calls that containers can make, limiting the attack surface. By using seccomp profiles, you can enforce what system calls a container is allowed to invoke.

apiVersion: v1

kind: Pod

metadata:

name: seccomp-pod

spec:

containers:

- name: seccomp-container

image: my-container

securityContext:

seccompProfile:

type: RuntimeDefault

In this example, the container uses the default seccomp profile, which restricts it from making potentially dangerous system calls.

b) AppArmor AppArmor is another Linux security module that provides mandatory access control (MAC) by restricting the capabilities of containers and their access to the underlying system. You can configure AppArmor profiles to restrict what a container can access or do.

apiVersion: v1

kind: Pod

metadata:

name: apparmor-pod

spec:

containers:

- name: apparmor-container

image: my-container

securityContext:

appArmorProfile: "runtime/default"

This example ensures that the container is constrained by the runtime/default AppArmor profile.

c) User Namespaces. Kubernetes can isolate the user namespace of containers, allowing containers to run with different user IDs and group IDs on the host. This helps to isolate containers and reduce the risks associated with running containers as root.

5. Best Practices for Application Security

- Run containers as non-root users: Always configure your containers to run as non-root users unless absolutely necessary.
- Limit the privileges of containers: Use SecurityContexts to limit capabilities and prevent privilege escalation.
- Use read-only root filesystem: Configure containers to use a read-only root filesystem to prevent modifications to container files.
- Use seccomp and AppArmor profiles: Implement seccomp and AppArmor to limit system calls and interactions with the host OS.
- Use RBAC for authorization: Implement Role-Based Access Control (RBAC) to restrict access to Kubernetes resources, ensuring only authorized entities can access sensitive resources.

- Audit security configurations: Regularly audit security settings like SecurityContexts, PSPs, and user privileges to ensure they adhere to the principle of least privilege.

6. Conclusion

In Kubernetes, application security is a vital part of ensuring that your workloads run in a controlled and secure environment. By leveraging tools like SecurityContexts, Linux capabilities, Pod Security Policies, seccomp, and AppArmor, you can minimize the risk of vulnerabilities and ensure your containers adhere to security best practices.

By implementing the principles of least privilege, restricting unnecessary capabilities, and utilizing advanced security features, you ensure that your Kubernetes applications are robust and secure, both inside the container and in the broader Kubernetes cluster.

Key takeaways:

- Always use SecurityContexts to define specific security configurations for Pods and containers.
- Limit container privileges using Linux capabilities and ensure Pod Security Policies are enforced.
- Leverage seccomp and AppArmor to restrict container interactions with the host.
- Continuously apply security best practices to safeguard your Kubernetes workloads.

Lab walkthrough: Understanding the need for controlling security in Containers

1. Run this command to see the possibilities of security context

```
kubectl explain pod.spec.containers.securityContext
```

2. Create a pod with privileged access. Put the following in a file called security-context-example.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-example
  namespace: test
spec:
  containers:
  - image: busybox
    name: busybox
    args:
    - sleep
    - "3600"
    securityContext:
      privileged: true
```

3. Apply the file

```
kubectl apply -f security-context-example.yaml
```

4. Open a shell in the pod:

```
kubectl exec -n test security-context-example -it -- /bin/sh
```

3. Mount the host's hard drive and read the kubelet kubeconfig file that is created during cluster initialization:

```
mkdir /hostfs
mount /dev/vda1 /hostfs
```

4. Now you can read the kubelet config file

```
cat /hostfs/kubelet/config.yaml
```

The kubeconfig file can be used to gain admin access to the Kubernetes cluster. Be careful with the containers you give privileged access to. You can control this with SecurityContext and Linux Capabilities.

Consider using a [PodSecurityPolicy](#) to lock down container privileges.

LAB Walkthrough 2: Set the security context for a Pod

To specify security settings for a Pod, include the `securityContext` field in the Pod specification. The `securityContext` field is a [PodSecurityContext](#) object. The security settings that you specify for a Pod apply to all Containers in the Pod. Here is a configuration file for a Pod that has a `securityContext` and an `emptyDir` volume:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: security-context-demo
```

```
spec:
```

```
  securityContext:
```

```
    runAsUser: 1000
```

```
    runAsGroup: 3000
```

```
    fsGroup: 2000
```

```
    supplementalGroups: [4000]
```

```
  volumes:
```

```
    - name: sec-ctx-vol
```

```
      emptyDir: {}
```

```
  containers:
```

```
    - name: sec-ctx-demo
```

```
      image: busybox:1.28
```

```
      command: [ "sh", "-c", "sleep 1h" ]
```


volumeMounts:

- name: sec-ctx-vol

mountPath: /data/demo

securityContext:

allowPrivilegeEscalation: false

- In the configuration file, the runAsUser field specifies that for any Containers in the Pod, all processes run with user ID 1000.
- The runAsGroup field specifies the primary group ID of 3000 for all processes within any containers of the Pod. If this field is omitted, the primary group ID of the containers will be root(0). Any files created will also be owned by user 1000 and group 3000 when runAsGroup is specified.
- Since fsGroup field is specified, all processes of the container are also part of the supplementary group ID 2000. The owner for volume /data/demo and any files created in that volume will be Group ID 2000.
- Additionally, when the supplementalGroups field is specified, all processes of the container are also part of the specified groups. If this field is omitted, it means empty.

1: Create the Pod:

```
kubectl apply -f security-context.yaml
```

2. Verify that the Pod's Container is running:

```
kubectl get pod security-context-demo
```

3. Get a shell to the running Container:

```
kubectl exec -it security-context-demo -- sh
```

4. In your shell, list the running processes:

```
ps
```

The output shows that the processes are running as user 1000, which is the value of runAsUser:

PID	USER	TIME	COMMAND
-----	------	------	---------

1	1000	0:00	sleep 1h
---	------	------	----------

6	1000	0:00	sh
---	------	------	----

...

5. In your shell, navigate to /data, and list the one directory:

```
cd /data
```

```
ls -l
```

The output shows that the /data/demo directory has group ID 2000, which is the value of fsGroup.

```
drwxrwsrwx 2 root 2000 4096 Jun  6 20:08 demo
```

6. In your shell, navigate to /data/demo, and create a file:

```
cd demo
```

```
echo hello > testfile
```

List the file in the /data/demo directory:

```
ls -l
```

The output shows that testfile has group ID 2000, which is the value of fsGroup.

```
-rw-r--r-- 1 1000 2000 6 Jun  6 20:08 testfile
```

7. Run the following command:

```
id
```

The output is similar to this:

```
uid=1000 gid=3000 groups=2000,3000,4000
```

From the output, you can see that gid is 3000 which is same as the runAsGroup field.

If the runAsGroup was omitted, the gid would remain as 0 (root) and the process will be able to interact with files that are owned by the root(0) group and groups that have the required group permissions for the root (0) group. You can also see that groups contains the group IDs which are specified by fsGroup and supplementalGroups, in addition to gid.

8. Exit your shell:

```
exit
```

Set the security context for a Container

To specify security settings for a Container, include the securityContext field in the Container manifest. The securityContext field is a [SecurityContext](#) object. Security settings that you specify for a Container apply only to the individual Container, and they override settings made at the Pod level when there is overlap. Container settings do not affect the Pod's Volumes.

Here is the configuration file for a Pod that has one Container. Both the Pod and the Container have a securityContext field:

security-context-2.yaml

apiVersion: v1

kind: Pod

metadata:

name: security-context-demo-2

spec:

securityContext:

runAsUser: 1000

containers:

- name: sec-ctx-demo-2

image: gcr.io/google-samples/hello-app:2.0

securityContext:

runAsUser: 2000

allowPrivilegeEscalation: false

1. Create the Pod:

```
kubectl apply -f https://k8s.io/examples/pods/security/security-context-2.yaml
```

2. Verify that the Pod's Container is running:

```
kubectl get pod security-context-demo-2
```

3. Get a shell into the running Container:

```
kubectl exec -it security-context-demo-2 -- sh
```

4. In your shell, list the running processes:

```
ps aux
```

The output shows that the processes are running as user 2000. This is the value of runAsUser specified for the Container. It overrides the value 1000 that is specified for the Pod.

```
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
2000        1  0.0  0.0  4336  764 ?        Ss   20:36   0:00 /bin/sh -c node server.js
2000        8  0.1  0.5 772124 22604 ?        Sl   20:36   0:00 node server.js
```

...

5. Exit your shell:

exit

Domain 3.5: Debugging in Kubernetes

Debugging is a critical skill for a Kubernetes Application Developer (CKAD). Kubernetes is a powerful platform for deploying and managing containerized applications, but it introduces complexity, especially when issues arise. Whether it's a misconfigured deployment, a failed Pod, or a networking issue, knowing how to debug and troubleshoot effectively is essential. This section will cover the tools and techniques you can use to debug applications and Kubernetes resources.

1. Key Debugging Concepts in Kubernetes

When debugging in Kubernetes, it's important to understand the following key concepts: Pod Lifecycle: Pods can go through various states like Pending, Running, Succeeded, or Failed. Understanding the lifecycle helps identify where a problem might be occurring.

1. Kubernetes Resources: Debugging Kubernetes resources (like Pods, Deployments, Services, and ConfigMaps) involves checking their configurations, status, and logs.
2. Cluster Context: Debugging often requires understanding the cluster context, including node status, kubelet logs, and network configuration.
3. Container Behavior: Problems with containers themselves (such as crashes, memory limits, or network timeouts) often manifest in the container logs.

2. Common Debugging Scenarios in Kubernetes

As an application developer, you'll face several common issues when working with Kubernetes. Here are some of the most frequent ones: Pods Not Starting: Pods that are stuck in the Pending or CrashLoopBackOff state. Service Unavailability: Pods are running but cannot be accessed via the Kubernetes Service. Configuration Issues: Incorrect ConfigMaps, Secrets, or environment variables affecting the behavior of the application. Resource Constraints: Pods or containers being terminated or restarted due to memory or CPU constraints.

3. Debugging Tools and Techniques

Kubernetes provides a variety of tools and commands to debug these issues. Here's a breakdown of the most useful debugging techniques:

1. kubectl describe

The kubectl describe command provides detailed information about Kubernetes resources. It's extremely useful for debugging issues with Pods, Services, Deployments, and more. Pods: To describe a Pod, use the following command:

```
kubectl describe pod <pod-name>
```

This will give you details about the Pod, including events, container status, resource usage, and any error messages.

Example output: If a Pod is stuck in Pending or CrashLoopBackOff, kubectl describe pod will often provide insight into what's going wrong (e.g., resource allocation, failed pull image, etc.).

2. kubectl logs

Logs are one of the most effective ways to understand what's happening inside a container. If a Pod is crashing or behaving unexpectedly, check its logs: Basic logs:

```
kubectll logs <pod-name> -c <container-name>
```

Stream logs in real-time:

```
kubectll logs -f <pod-name> -c <container-name>
```

Previous logs: If the container crashed and was restarted, you can view logs from the previous instance:

```
kubectll logs <pod-name> -c <container-name> --previous
```

Logs often provide the most direct insight into why a container is failing, such as uncaught exceptions, resource issues, or missing dependencies.

3. kubectll get

The kubectll get command is used to retrieve information about Kubernetes resources. It can be helpful to view the current state of resources and identify any issues. Pods: To get information about the status of Pods in a namespace:

```
kubectll get pods -n <namespace>
```

Deployment: To check the status of a deployment:

```
kubectll get deployment <deployment-name> -n <namespace>
```

Services: To check the Services running in the cluster:

```
kubectll get svc -n <namespace>
```

The output will show the current state, such as whether the Pods are running, and if there are any issues with resource allocation.

4. kubectll top

The kubectll top command is used to monitor resource usage (CPU and memory) for nodes and Pods. It can help identify if your application is running out of resources. Monitor nodes:

```
kubectll top nodes
```

Monitor Pods:

```
kubectll top pods -n <namespace>
```

If a Pod is experiencing memory or CPU issues, it may be terminated or fail to start. `kubectl top` helps identify these problems.

5. `kubectl port-forward`

If you need to debug an application from inside a Pod, the `kubectl port-forward` command allows you to access an application running inside a Kubernetes cluster without exposing it externally. Example command:

```
kubectl port-forward pod/<pod-name> 8080:80
```

This forwards port 8080 on your local machine to port 80 inside the container. This is useful for debugging an application's behavior from your local machine without exposing it to the public internet.

6. `kubectl exec`

The `kubectl exec` command allows you to execute commands directly inside a running container. This is useful for inspecting the state of an application or running debugging tools in the container. Example command:

```
kubectl exec -it <pod-name> -c <container-name> -- /bin/bash
```

This opens an interactive terminal session inside the container, allowing you to run commands such as `ps`, `top`, `curl`, etc., to investigate the issue.

7. Checking Events

Kubernetes events provide insights into what's happening in the cluster, including Pod creation, failure, and resource allocation issues. You can view events using:

```
kubectl get events -n <namespace>
```

Events help identify why a Pod might be stuck in the Pending state, why it was evicted, or why a deployment failed to rollout.

4. Troubleshooting Common Issues

Here are some common issues and how to debug them:

1. Pod Stuck in Pending

Check the Pod's events: Often, Pods stuck in the Pending state are waiting for resources (like CPU or memory). Use

```
kubectl describe pod <pod-name>
```

to see if there are any warnings about resource availability.

Check resource requests/limits: If your Pod is requesting more resources than are available, it may not be scheduled. Check if the Pod has resource requests/limits set appropriately.

2. Pod Crashes or Restarts (CrashLoopBackOff)

View logs:

Use `kubectl logs <pod-name>`

to see what's happening in the container. Look for errors like missing files, connection issues, or configuration problems.

- Check for resource issues: Make sure the container isn't being killed due to resource limits (memory/CPU). Use `kubectl top pods` to monitor usage.
- Check readiness/liveness probes: A misconfigured probe can cause Kubernetes to restart the Pod frequently. Review the health check configurations.

3. Service Not Accessible

- Verify Service: Use `kubectl describe svc` to ensure the Service is correctly pointing to the Pods and has the correct ports.
- Check Network Policies: Ensure there are no network policies blocking access to the service.
- Check Pod Logs: If the application inside the Pod is failing to respond to requests, check the logs of the container to ensure the application is running and listening on the expected port.

4. Misconfigured Environment Variables or ConfigMaps

- Inspect ConfigMap/Secret: Use `kubectl describe configmap` or `kubectl describe secret` to ensure that the correct environment variables are being set.
- Check Pods' Environment: Use `kubectl exec` to inspect the environment inside the container to ensure that variables are correctly passed.

5. Debugging Network Issues

Networking problems often occur when Pods are unable to communicate with each other, a service, or the outside world. Common network issues include DNS resolution problems, service misconfiguration, or blocked ports.

- Test Connectivity: Use `kubectl exec` to ping other Pods or services to verify connectivity.
- Check Network Policies: Ensure that Kubernetes network policies are not unintentionally blocking traffic.
- Check DNS Resolution: If Pods cannot resolve domain names, it might indicate a DNS issue. Check the CoreDNS logs using `kubectl logs -n kube-system`.

6. Conclusion

Debugging is a crucial skill for a Kubernetes Application Developer (CKAD). Kubernetes provides several powerful tools and commands like

`kubectl describe`, `kubectl logs`, `kubectl exec`, and `kubectl top`

to help you troubleshoot and identify issues with Pods, containers, and other resources.

By familiarizing yourself with these tools and techniques, you will be better equipped to identify, diagnose, and resolve issues in your Kubernetes-based applications, ensuring smooth operation in a cloud-native environment.

Domain 4.7: Understand Service Accounts

In Kubernetes, Service Accounts are used to provide an identity for processes that run in Pods. When applications in Pods need to interact with the Kubernetes API or other resources within the cluster, they use a ServiceAccount to authenticate and authorize themselves. Each ServiceAccount is associated with a set of API access credentials, typically tokens, that the system uses to authenticate requests made by Pods. Understanding Service Accounts is crucial for building secure Kubernetes applications, as they define how Pods interact with the Kubernetes API and other resources in the cluster.

In this section, we will cover the following topics:

- What Service Accounts are and why they are important.
- How to create and configure Service Accounts.
- How to associate Service Accounts with Pods.
- Best practices for using Service Accounts securely.

1. What is a Kubernetes Service Account?

A ServiceAccount is a Kubernetes resource that provides an identity for a Pod and is used to authenticate and authorize API requests on behalf of the Pod. It's essentially a "service user" that can be assigned specific permissions (via RBAC policies) to interact with other resources within the Kubernetes cluster.

Key Points:

- Service Accounts are tied to namespaces. Each namespace can have its own set of Service Accounts.
- A ServiceAccount is associated with an automatically generated ServiceAccount Token (JWT) that is used for authentication.
- Service Accounts can be configured to interact with the Kubernetes API, make calls to other services, or access secrets.

2. Why are Service Accounts Important?

Service Accounts provide a way to manage the security and identity of applications running in your Kubernetes clusters. They help in:

- Access Control: Service Accounts are used with Role-Based Access Control (RBAC) to define what a Pod can and cannot do in the Kubernetes environment.
- Isolation and Least Privilege: By using distinct Service Accounts for different types of workloads, you can minimize the risk of exposing sensitive data or giving unnecessary permissions.
- Automation: Service Accounts are key for applications that need to automatically interact with the Kubernetes API or other resources in the cluster, such as creating Pods, accessing secrets, etc.

3. Default Service Account

Every Kubernetes namespace automatically has a default ServiceAccount. When you create a Pod and don't explicitly specify a ServiceAccount, Kubernetes automatically assigns the default ServiceAccount of the namespace to the Pod.

Example of Using Default Service Account

If you create a Pod without specifying a ServiceAccount, it will use the default one:

apiVersion: v1

kind: Pod

metadata:

name: default-service-account-pod

spec:

containers:

- name: my-container

image: my-image

This Pod will automatically use the default ServiceAccount from the same namespace.

4. Creating and Using Service Accounts

a) Creating a Service Account

You can create a Service Account using a simple YAML manifest or the kubectl command. Here's an example of a ServiceAccount YAML manifest:

apiVersion: v1

kind: ServiceAccount

metadata:

name: my-service-account

To create the ServiceAccount:

kubectl apply -f service-account.yaml

Alternatively, you can create a ServiceAccount directly via kubectl:

kubectl create serviceaccount my-service-account

b) Associating a Service Account with a Pod

Once a ServiceAccount is created, you can associate it with a Pod by specifying it in the Pod's serviceAccountName field:

apiVersion: v1

kind: Pod

metadata:

```
name: pod-with-service-account
```

```
spec:
```

```
serviceAccountName: my-service-account
```

```
containers:
```

```
- name: my-container
```

```
image: my-image
```

This Pod will use the my-service-account ServiceAccount for authentication and authorization when interacting with the Kubernetes API. If you don't explicitly set the serviceAccountName, the Pod will use the default ServiceAccount in the namespace.

5. Service Account Tokens

When a Pod uses a ServiceAccount, Kubernetes automatically creates a ServiceAccount Token. This token is used by the Pod to authenticate with the Kubernetes API.

Kubernetes automatically mounts the token as a file inside the Pod, typically at /var/run/secrets/kubernetes.io/serviceaccount/token.

The token is in JWT (JSON Web Token) format and can be used for authenticating API requests.

Example: Accessing the ServiceAccount Token Inside a Pod

apiVersion: v1

```
kind: Pod
```

```
metadata:
```

```
name: pod-with-token
```

```
spec:
```

```
serviceAccountName: my-service-account
```

```
containers:
```

```
- name: my-container
```

```
image: my-image
```

```
command: ["cat", "/var/run/secrets/kubernetes.io/serviceaccount/token"]
```

This Pod will print the ServiceAccount token when it is run.

6. Role-Based Access Control (RBAC) with Service Accounts

Service Accounts are often used in combination with Role-Based Access Control (RBAC) to define what a ServiceAccount (and, by extension, the Pod using that ServiceAccount) can do in the cluster.

a) Create a Role

Here's an example of creating a Role that allows access to Pods within a namespace:

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

metadata:

namespace: default

name: pod-reader

rules:

- apiGroups: [""]

resources: ["pods"]

verbs: ["get", "list"]

This Role gives the ServiceAccount permission to get and list Pods in the default namespace.

b) Create a RoleBinding

A RoleBinding associates the Role with a ServiceAccount: apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

name: pod-reader-binding

namespace: default

subjects:

- kind: ServiceAccount

name: my-service-account

namespace: default

roleRef:

kind: Role

name: pod-reader

apiGroup: rbac.authorization.k8s.io

This binding allows the my-service-account ServiceAccount to read Pods in the default namespace.

7. Service Account Best Practices

1. Use the Least Privilege Principle: Grant only the minimum permissions required for the ServiceAccount to perform its tasks. This minimizes the risk of exposure.
2. Use Separate Service Accounts for Different Workloads: Avoid using the default ServiceAccount for every Pod. Create distinct ServiceAccounts for different types of workloads or applications to reduce the blast radius of compromised credentials.
3. Avoid Sharing Service Accounts Across Pods: When multiple Pods need different levels of access, create separate ServiceAccounts with the appropriate permissions for each.
4. Audit Service Account Usage: Regularly audit the permissions and roles granted to ServiceAccounts, and ensure they are not over-privileged.
5. Secure the Service Account Tokens: Ensure tokens are used securely, and access to the file paths where tokens are stored (e.g., /var/run/secrets/kubernetes.io/serviceaccount/) is restricted.
6. Rotate Service Account Tokens: Kubernetes automatically handles the expiration of tokens, but for critical workloads, consider integrating with external tools for more fine-grained control and rotation.

8. Deleting a Service Account

To delete a ServiceAccount:

```
kubectl delete serviceaccount my-service-account
```

Make sure to check if any Pods are using the ServiceAccount before deletion, as removing it may affect those Pods' ability to interact with the Kubernetes API.

9 . Conclusion

Service Accounts are a vital component in Kubernetes security, providing a way to authenticate and authorize Pods and workloads within a cluster. By associating Service Accounts with appropriate RBAC roles, you can enforce strict access control policies and reduce the risk of over-privileged access.

In this section, you have learned:

- How to create and use Service Accounts.
- How to associate Service Accounts with Pods.
- The role of RBAC in controlling access to Kubernetes resources via Service Accounts.
- Best practices for securing and managing Service Accounts in your cluster.

By following best practices and using Service Accounts appropriately, you can secure your Kubernetes applications and ensure that they interact with the Kubernetes API and other resources in a controlled and authorized manner.

Domain 3.1: Understand API Deprecations

As Kubernetes evolves, the APIs and features it provides can change over time. API deprecations are an essential aspect of this evolution. Understanding how Kubernetes manages deprecations and how to handle them in your application development is crucial for ensuring that your workloads remain stable and compatible with future versions of Kubernetes.

In this section, we'll explore what API deprecations are, why they occur, and how to handle them effectively as a Kubernetes Application Developer (CKAD).

1. What Are API Deprecations?

In Kubernetes, API deprecation refers to the process of marking an API or API version as outdated and signaling that it will be removed in a future release. The deprecation process is Kubernetes' way of ensuring backward compatibility while allowing users time to update their applications and configurations before older versions of APIs are removed.

Why Do API Deprecations Happen?

- **Improved Features:** New versions of APIs may provide more efficient, secure, or feature-rich implementations.
- **Code Cleanup:** Over time, certain API versions may no longer be needed as Kubernetes evolves, and these old versions are deprecated.
- **Enhanced Performance & Usability:** The Kubernetes team may optimize APIs, change field names, or restructure resources to improve performance, usability, and flexibility. When an API is deprecated, it is typically still available for some time but may show warnings about its deprecation. Eventually, it will be removed in a future release (often a major version).

2. How Kubernetes Marks APIs as Deprecated

Kubernetes marks an API as deprecated in the following ways:

- **Deprecation Warnings in Logs:** When using a deprecated API, Kubernetes emits warning messages indicating the API is deprecated and should be updated.
- **Documentation:** Kubernetes documentation provides clear guidance on which API versions are deprecated and what their replacements are.
- **API Versioning:** Kubernetes has multiple versions of APIs to ensure backward compatibility.

Deprecated APIs are often maintained for a while in a version before they are completely removed.

For example, if you are using an older version of an API, such as `extensions/v1beta1` for resources like Deployments, you might receive a warning about deprecation in the logs:

```
W0311 16:10:23.123456      1 warnings.go:67] extensions/v1beta1 deployments is deprecated in v1.19.0+, use apps/v1 instead.
```

3. Common Kubernetes APIs That Are Deprecated

Over the years, Kubernetes has deprecated several APIs. Some of the most notable deprecated APIs include:

extensions/v1beta1 -> apps/v1 (Deployments, ReplicaSets, etc.):

The extensions/v1beta1 API for resources like Deployments and ReplicaSets was deprecated in favor of the new API group apps/v1. Example:

```
apiVersion: extensions/v1beta1 # Deprecated
```

```
kind: Deployment
```

extensions/v1beta1 -> networking.k8s.io/v1 (Ingress):

The extensions/v1beta1 API for Ingress resources was also deprecated in favor of networking.k8s.io/v1.

Example:

```
apiVersion: networking.k8s.io/v1 # New
```

```
kind: Ingress
```

apiextensions.k8s.io/v1beta1 -> apiextensions.k8s.io/v1 (CustomResourceDefinitions):

The v1beta1 version of CustomResourceDefinitions (CRDs) was deprecated in favor of v1.

batch/v1beta1 -> batch/v1 (CronJobs):

The batch/v1beta1 version of CronJobs was deprecated in favor of batch/v1.

4. How to Handle API Deprecations

As a Kubernetes Application Developer (CKAD), it's important to be proactive about handling API deprecations. Here's how you can manage deprecated APIs:

1. Stay Updated with Kubernetes Release Notes

Kubernetes releases new versions regularly (every 3 months), and each release contains a changelog that lists deprecated APIs and their replacements. Staying on top of these changes will allow you to plan your upgrades and refactor your application accordingly. You can find release notes on the official Kubernetes GitHub page and on Kubernetes documentation.

2. Update Your YAML Manifests

Whenever you encounter deprecation warnings in your Kubernetes environment, you should update your YAML files to use the newer API versions. For example, if you're using extensions/v1beta1 for Deployments, you should switch to apps/v1.


```
apiVersion: apps/v1 # New API version
```

```
kind: Deployment
```

```
metadata:
```

```
  name: my-app
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      app: my-app
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: my-app
```

```
    spec:
```

```
      containers:
```

```
        - name: my-app
```

```
          image: my-app:latest
```

3. Use kubectl to Identify Deprecated APIs

You can use the kubectl tool to check for deprecated API versions in your cluster. The kubectl deprecations plugin is available to detect deprecated APIs in your cluster resources:

```
kubectl deprecations
```

This tool will give you an overview of the deprecated APIs being used in your cluster and provide guidance on how to update them.

4. Use API Version Migrations

If the deprecated API is still functional, Kubernetes will generally provide a migration path. Kubernetes' API reference documentation and release notes will help you understand how to migrate from deprecated versions to newer, stable ones. For example:

Migrate from extensions/v1beta1 to apps/v1 for Deployments.

Migrate from `extensions/v1beta1` to `networking.k8s.io/v1` for Ingress.

5. Test Changes in Staging Environments

Before deploying changes that involve deprecated API migrations to production, test your application thoroughly in a staging environment. This ensures that the new API versions don't break your application and that the transition is smooth.

6. Leverage Kubernetes Features for Compatibility

Kubernetes supports multi-version compatibility for a certain period. This means you can run multiple versions of APIs for a while (e.g., `apps/v1` and `extensions/v1beta1`) in the same cluster. However, you should plan to update deprecated APIs before the versions are removed completely in future Kubernetes releases.

5. Deprecation Timeline

Kubernetes provides a deprecation timeline in its release notes. Kubernetes uses a two-year deprecation policy:

- A feature marked as deprecated will be supported for two more major versions (around 6 months each).
- After the deprecation period ends, the feature is removed.

For instance:

Kubernetes 1.14 marked `extensions/v1beta1` APIs (like Deployments and Ingress) as deprecated.

- They were fully removed in Kubernetes 1.22.
- The deprecation timeline is announced in the release notes, so it's important to stay informed about upcoming changes and plan your application updates accordingly.

6. Benefits of Understanding API Deprecations

- **Future-proofing:** By addressing deprecated APIs early, you ensure that your applications remain compatible with future Kubernetes versions.
- **Security:** Newer API versions often come with security improvements, bug fixes, and better performance.
- **Stability:** By updating deprecated APIs, you avoid the risk of them being removed in future versions, which could break your workloads.
- **Improved Maintainability:** Using up-to-date API versions ensures that your Kubernetes manifests are aligned with the current best practices and Kubernetes standards.

7. Conclusion

As a Kubernetes Application Developer (CKAD), staying on top of API deprecations is a critical skill to ensure that your applications remain compatible with newer versions of Kubernetes. By:

1. Regularly reviewing release notes for deprecated APIs.
2. Updating your YAML configurations to use current APIs.
3. Testing changes in non-production environments.

You can ensure your Kubernetes applications are future-proof and continue to work seamlessly as Kubernetes evolves. Pro tip: Utilize tools like `kubectl deprecations` and check Kubernetes' release notes to track deprecations and plan your updates accordingly.

Domain 3.4: Use Built-in CLI Tools to Monitor Kubernetes Applications

As a Kubernetes Application Developer (CKAD), it's essential to be able to monitor the applications you deploy on Kubernetes. Monitoring helps you track the health, performance, and availability of applications, and identify any issues in real time. Kubernetes provides built-in CLI tools that can help you monitor applications and gain insight into the status of your resources. In this section, we'll focus on some of the most commonly used Kubernetes CLI tools and commands for monitoring your applications and how to use them effectively.

1. kubectl - The Kubernetes Command-Line Interface

The kubectl CLI is the primary tool for interacting with Kubernetes clusters. It can be used to inspect and monitor various components of the cluster and applications running in it.

2. Common kubectl Commands for Monitoring Kubernetes Applications

1. Checking Pod Status

You can use kubectl to get information about your Pods, such as their current state, events, and resource usage. This is useful for monitoring the health and status of the applications within your Pods. View all Pods in a namespace:

```
kubectl get pods
```

Use

```
kubectl get pods -n <namespace>
```

to get Pods from a specific namespace.

Use `kubectl get pods -o wide`

to show additional details like the node the pod is running on. View detailed information about a specific Pod:

```
kubectl describe pod <pod-name>
```

This command gives you detailed information about the Pod, including events, container status, resource usage, and more.

View the logs of a specific Pod:

```
kubectl logs <pod-name>
```

Use `-f` to stream the logs in real time:

```
kubectl logs -f <pod-name>
```

View logs for all containers in a Pod (if there are multiple containers):
`kubectll logs <pod-name> -c <container-name>`

2. Monitoring Deployments

A Deployment is a higher-level object in Kubernetes used to manage the lifecycle of Pods. You can use `kubectll` to monitor the status of Deployments. View all Deployments:
`kubectll get deployments`

Get detailed information about a specific Deployment:
`kubectll describe deployment <deployment-name>`

This provides insights into the Deployment's status, including the number of replicas, pods being used, and any ongoing issues. Check the rollout status of a Deployment:
`kubectll rollout status deployment/<deployment-name>`

This command helps you track the progress of a rolling update and verify that the deployment is working as expected. Undo a Deployment rollout:
If something goes wrong, you can rollback to the previous version:
`kubectll rollout undo deployment/<deployment-name>`

3. Monitoring Services

A Service in Kubernetes exposes your application to the network. Monitoring Services helps ensure that your application is accessible and traffic is being routed correctly. View all Services in a namespace:
`kubectll get services`

Get detailed information about a specific Service:
`kubectll describe service <service-name>`

This will show details about the service, such as the IP address, ports, and which Pods are being targeted.

4. Monitoring Namespaces

Namespaces are a way to partition resources in a Kubernetes cluster. Monitoring namespaces helps you get a clearer understanding of resource allocation and application isolation. List all namespaces:

```
kubectl get namespaces
```

Get resources (Pods, Services, Deployments) within a specific namespace:

```
kubectl get all -n <namespace>
```

5. Monitoring Resource Usage

Kubernetes also provides resource metrics, such as CPU and memory usage, to help you monitor how efficiently your applications are running. View resource usage for Pods:

To use this command, you need to have the metrics-server installed in your cluster.

```
kubectl top pod
```

This command shows the current CPU and memory usage for all Pods in the default namespace. You can also specify a namespace:

```
kubectl top pod -n <namespace>
```

View resource usage for Nodes:

```
kubectl top node
```

This command shows the CPU and memory usage for each node in the cluster.

6. Monitoring Events

Kubernetes logs events that provide insights into what's happening in the cluster, such as Pod failures, resource allocation issues, or deployment status changes. View all events in a namespace:

```
kubectl get events -n <namespace>
```

This command shows all events in a specific namespace. Use it to track errors, warnings, or changes to your application or resources.

View events in real time:

```
kubectl get events --watch
```

This allows you to stream events as they occur in your cluster, which is helpful for troubleshooting live issues.

3. Monitoring with kubectl Plugins

Kubernetes supports the use of plugins that can extend the functionality of kubectl. Some plugins are designed specifically for monitoring and troubleshooting.

1. kube-ps1: Provides Kubernetes context information in your shell prompt.
 2. kubectl-gotmpl: Allows you to use Go templating for more flexible queries and outputs.
 3. kubectl-debug: Provides an easy way to debug running Pods by creating a temporary container.
- You can install these plugins to enhance your Kubernetes monitoring experience.

4. Troubleshooting Application Failures with kubectl

When things go wrong, you can use kubectl to troubleshoot the issue. Inspect the state of Pods:
kubectl get pods -o wide

This command shows which node the Pods are running on and helps you diagnose if there is a resource issue (e.g., resource limits, node failures).

Describe the Pod to get detailed information about the error:

kubectl describe pod <pod-name>

This provides detailed event logs and error messages for the container and helps you identify why the Pod might not be running properly.

Access a shell inside a running Pod (for debugging purposes):

kubectl exec -it <pod-name> -- /bin/sh

This allows you to interact with the application running inside the container and inspect logs, files, or processes.

5. Leveraging Labels and Selectors Kubernetes allows you to use labels and selectors to filter resources in your cluster. Labels can be added to Pods, Services, and other resources to group them based on common attributes. Get Pods with a specific label:

kubectl get pods -l app=<label-name>

View resources by label in a specific namespace:

kubectl get all -l app=<label-name> -n <namespace>

Using labels and selectors is helpful when you need to monitor or troubleshoot a subset of resources in a large cluster.

6. Conclusion

As a Kubernetes Application Developer (CKAD), being proficient with kubectl is essential for monitoring the health, performance, and availability of your applications in a Kubernetes cluster. By using the built-in CLI tools, you can quickly gain insight into the state of your Pods, Deployments, Services, and other resources, allowing you to troubleshoot and optimize your applications effectively.

Some key commands to remember:

1. `kubectl get` — View resources (Pods, Deployments, Services, etc.).
2. `kubectl describe` — Get detailed resource information and events.
3. `kubectl logs` — View application logs.
4. `kubectl top` — Monitor resource usage.
5. `kubectl exec` — Access running Pods for debugging.

Mastering these tools will empower you to ensure your applications are running smoothly in your Kubernetes environment, enabling you to detect and resolve issues quickly.

Domain 3.4: Utilize Container Logs

As a Kubernetes Application Developer (CKAD), one of your core responsibilities is ensuring the health and performance of applications running in Kubernetes. Logs are a critical part of this process, as they provide detailed insights into what's happening inside containers and help you troubleshoot issues. Kubernetes provides a robust system for capturing and managing container logs, which can be accessed through various means using `kubectl` and integrated tools. In this section, we'll discuss how to effectively utilize container logs in Kubernetes for monitoring, debugging, and troubleshooting applications.

1. Basics of Container Logging in Kubernetes

Kubernetes stores logs from containers in a centralized location on the node filesystem. By default, logs are captured from the standard output (`stdout`) and standard error (`stderr`) streams of the containers running in your Pods. These logs are then made accessible to you through the `kubectl` CLI.

- **Pod Logs:** Logs are collected per Pod and per container within the Pod.
- **Log Location:** On each node, container logs are typically stored in the `/var/log/containers` directory. However, accessing logs through `kubectl` is the recommended way.

2. Accessing Container Logs Using `kubectl`

1. View Logs of a Specific Pod

To view the logs of a container within a specific Pod, use the following `kubectl` command:
`kubectl logs <pod-name>`

This will retrieve the logs from the first container in the Pod. Specify the container name (if the Pod has multiple containers):

```
kubectl logs <pod-name> -c <container-name>
```

View logs for a specific namespace:

```
kubectl logs <pod-name> -n <namespace>
```

2. View Previous Logs

If a container has crashed and restarted, you can view the logs from the previous instance of the container by using the `--previous` flag:

```
kubectl logs <pod-name> -c <container-name> --previous
```

This is useful for debugging container crashes or other issues that occur before the container restarts.

3. Stream Logs in Real-Time

To monitor logs in real time (i.e., stream the logs as they are generated), use the `-f` (follow) option:
`kubectl logs -f <pod-name> -c <container-name>`

This command keeps the connection open and continuously streams logs as they are generated, which is invaluable when debugging live issues.

3. Viewing Logs for All Containers in a Pod

If a Pod has multiple containers, you can view the logs for all containers simultaneously by using the following command:

```
kubectl logs <pod-name> --all-containers=true
```

This will display the logs for every container within the specified Pod.

4. Retrieve Logs from All Pods in a Namespace

If you need to monitor logs for multiple Pods in a namespace, you can run:

```
kubectl logs -n <namespace> --selector=<label-selector> --all-containers=true
```

The `--selector=` helps filter the Pods based on labels, which can be useful if you only want to view logs for a specific set of Pods (e.g., Pods with a particular application label).

The `--all-containers=true` flag ensures logs from all containers in the selected Pods are displayed.

5. Searching and Filtering Logs

If you're looking for specific information in your logs, you can pipe the output of `kubectl logs` into common Unix commands like `grep` to search for patterns. For example, to find all occurrences of the string "error" in the logs:

```
kubectl logs <pod-name> -c <container-name> | grep "error"
```

This can help you quickly identify issues and narrow down your troubleshooting.

6. Log Aggregation in Kubernetes

While `kubectl` provides access to logs, for large-scale Kubernetes environments, it's recommended to use a log aggregation solution. These solutions aggregate and centralize logs from all your Pods, Nodes, and containers, making it easier to search, analyze, and monitor logs across your entire Kubernetes cluster. Some common log aggregation tools include:

1. ELK Stack (Elasticsearch, Logstash, and Kibana)
2. Fluentd
3. Prometheus & Grafana (used for metrics but can also handle logs with integrations)
4. Loki (part of the Grafana ecosystem)

By using these tools, logs are collected, indexed, and stored in a centralized location, allowing you to search and visualize logs in real-time through dashboards.

7. Implementing Logging Best Practices

To make your logs more useful and easier to manage, you can adopt the following logging best practices:

1. Structured Logging Ensure that your applications generate logs in a structured format, such as JSON. Structured logs are easier to parse and analyze, especially when using log aggregation tools. Example JSON log entry:

```
{
  "level": "error",
  "message": "Failed to connect to the database",
  "timestamp": "2023-01-10T10:00:00Z",
  "pod_name": "my-app-12345",
  "container_name": "my-app-container"
}
```

2. Use Log Levels Use log levels such as INFO, WARN, ERROR, and DEBUG to classify the severity of logs. This makes it easier to filter and prioritize issues when viewing logs.
3. Keep Logs Short and Relevant While it's important to log enough information to understand what's going on, avoid over-logging. Logs should be concise, relevant, and avoid redundancy.
4. Set Log Rotation In large-scale Kubernetes environments, logs can grow quickly. Ensure that log rotation is configured so logs don't fill up disk space. This can be managed by tools like logrotate on the node or within the logging solution you use (such as Fluentd or the ELK stack).
5. Manage and Monitor Log Retention Implement a retention policy for logs to avoid the accumulation of large amounts of data. You can define how long logs should be stored, and ensure they're archived or deleted after a certain period.

8. Troubleshooting with Container Logs

Logs are your first line of defense when troubleshooting container-related issues. Here's how you can troubleshoot common problems using container logs:

1. Pod Crash Looping If a container within a Pod is in a crash loop, you can inspect the logs to find out why the application is failing.

```
kubectl logs <pod-name> -c <container-name> --previous
```

Look for error messages or stack traces that indicate the cause of failure.

2. Application Behavior If your application isn't behaving as expected, logs can provide clues. Use `kubectl logs -f` to stream the logs and monitor application activity over time.
3. Networking Issues If your application is having networking issues (e.g., unable to connect to other services or databases), checking the logs can reveal errors related to network connectivity.
4. Resource Constraints If your application is failing due to resource constraints (e.g., CPU or memory limits), you can examine the logs to find out if it's hitting limits or being throttled.
5. Conclusion As a Kubernetes Application Developer (CKAD), utilizing container logs effectively is crucial for ensuring the health of your applications. Kubernetes makes it easy to access logs with the `kubectl logs` command, but for large-scale environments, implementing a log aggregation and analysis tool is essential for efficient monitoring.

Some best practices for utilizing container logs:

1. Use structured logging and log levels to make logs easier to parse.
2. Stream logs in real-time to monitor live events.
3. Use log aggregation tools to centralize and search logs efficiently.
4. Monitor logs for patterns and common issues such as crashes or resource limits.
5. Mastering the use of container logs will help you quickly identify and resolve issues, leading to more stable and reliable applications on Kubernetes.

Domain 3.5: Debugging in Kubernetes

Debugging is a critical skill for a Kubernetes Application Developer (CKAD). Kubernetes is a powerful platform for deploying and managing containerized applications, but it introduces complexity, especially when issues arise. Whether it's a misconfigured deployment, a failed Pod, or a networking issue, knowing how to debug and troubleshoot effectively is essential. This section will cover the tools and techniques you can use to debug applications and Kubernetes resources.

1. Key Debugging Concepts in Kubernetes

When debugging in Kubernetes, it's important to understand the following key concepts: Pod Lifecycle: Pods can go through various states like Pending, Running, Succeeded, or Failed. Understanding the lifecycle helps identify where a problem might be occurring.

1. Kubernetes Resources: Debugging Kubernetes resources (like Pods, Deployments, Services, and ConfigMaps) involves checking their configurations, status, and logs.
2. Cluster Context: Debugging often requires understanding the cluster context, including node status, kubelet logs, and network configuration.
3. Container Behavior: Problems with containers themselves (such as crashes, memory limits, or network timeouts) often manifest in the container logs.

2. Common Debugging Scenarios in Kubernetes

As an application developer, you'll face several common issues when working with Kubernetes. Here are some of the most frequent ones: Pods Not Starting: Pods that are stuck in the Pending or CrashLoopBackOff state. Service Unavailability: Pods are running but cannot be accessed via the Kubernetes Service. Configuration Issues: Incorrect ConfigMaps, Secrets, or environment variables affecting the behavior of the application. Resource Constraints: Pods or containers being terminated or restarted due to memory or CPU constraints.

3. Debugging Tools and Techniques

Kubernetes provides a variety of tools and commands to debug these issues. Here's a breakdown of the most useful debugging techniques:

1. kubectl describe

The kubectl describe command provides detailed information about Kubernetes resources. It's extremely useful for debugging issues with Pods, Services, Deployments, and more. Pods: To describe a Pod, use the following command:

```
kubectl describe pod <pod-name>
```

This will give you details about the Pod, including events, container status, resource usage, and any error messages.

Example output: If a Pod is stuck in Pending or CrashLoopBackOff, kubectl describe pod will often provide insight into what's going wrong (e.g., resource allocation, failed pull image, etc.).

2. kubectl logs

Logs are one of the most effective ways to understand what's happening inside a container. If a Pod is crashing or behaving unexpectedly, check its logs: Basic logs:

```
kubectl logs <pod-name> -c <container-name>
```

Stream logs in real-time:

```
kubectl logs -f <pod-name> -c <container-name>
```

Previous logs: If the container crashed and was restarted, you can view logs from the previous instance:

```
kubectl logs <pod-name> -c <container-name> --previous
```

Logs often provide the most direct insight into why a container is failing, such as uncaught exceptions, resource issues, or missing dependencies.

3. kubectl get

The kubectl get command is used to retrieve information about Kubernetes resources. It can be helpful to view the current state of resources and identify any issues. Pods: To get information about the status of Pods in a namespace:

```
kubectl get pods -n <namespace>
```

Deployment: To check the status of a deployment:

```
kubectl get deployment <deployment-name> -n <namespace>
```

Services: To check the Services running in the cluster:

```
kubectl get svc -n <namespace>
```

The output will show the current state, such as whether the Pods are running, and if there are any issues with resource allocation.

4. kubectl top

The kubectl top command is used to monitor resource usage (CPU and memory) for nodes and Pods. It can help identify if your application is running out of resources. Monitor nodes:

```
kubectl top nodes
```

Monitor Pods:

```
kubectl top pods -n <namespace>
```

If a Pod is experiencing memory or CPU issues, it may be terminated or fail to start. kubectl top helps identify these problems.

5. kubectl port-forward

If you need to debug an application from inside a Pod, the `kubectl port-forward` command allows you to access an application running inside a Kubernetes cluster without exposing it externally. Example command:

```
kubectl port-forward pod/<pod-name> 8080:80
```

This forwards port 8080 on your local machine to port 80 inside the container. This is useful for debugging an application's behavior from your local machine without exposing it to the public internet.

6. kubectl exec

The `kubectl exec` command allows you to execute commands directly inside a running container. This is useful for inspecting the state of an application or running debugging tools in the container. Example command:

```
kubectl exec -it <pod-name> -c <container-name> -- /bin/bash
```

This opens an interactive terminal session inside the container, allowing you to run commands such as `ps`, `top`, `curl`, etc., to investigate the issue.

7. Checking Events

Kubernetes events provide insights into what's happening in the cluster, including Pod creation, failure, and resource allocation issues. You can view events using:

```
kubectl get events -n <namespace>
```

Events help identify why a Pod might be stuck in the Pending state, why it was evicted, or why a deployment failed to rollout.

4. Troubleshooting Common Issues

Here are some common issues and how to debug them:

1. Pod Stuck in Pending

Check the Pod's events: Often, Pods stuck in the Pending state are waiting for resources (like CPU or memory). Use

```
kubectl describe pod <pod-name>
```

to see if there are any warnings about resource availability.

Check resource requests/limits: If your Pod is requesting more resources than are available, it may not be scheduled. Check if the Pod has resource requests/limits set appropriately.

2. Pod Crashes or Restarts (CrashLoopBackOff)

View logs:
Use `kubectl logs <pod-name>`

to see what's happening in the container. Look for errors like missing files, connection issues, or configuration problems.

- Check for resource issues: Make sure the container isn't being killed due to resource limits (memory/CPU). Use `kubectl top pods` to monitor usage.
- Check readiness/liveness probes: A misconfigured probe can cause Kubernetes to restart the Pod frequently. Review the health check configurations.

3. Service Not Accessible

- Verify Service: Use `kubectl describe svc` to ensure the Service is correctly pointing to the Pods and has the correct ports.
- Check Network Policies: Ensure there are no network policies blocking access to the service.
- Check Pod Logs: If the application inside the Pod is failing to respond to requests, check the logs of the container to ensure the application is running and listening on the expected port.

4. Misconfigured Environment Variables or ConfigMaps

- Inspect ConfigMap/Secret: Use `kubectl describe configmap` or `kubectl describe secret` to ensure that the correct environment variables are being set.
- Check Pods' Environment: Use `kubectl exec` to inspect the environment inside the container to ensure that variables are correctly passed.

5. Debugging Network Issues

Networking problems often occur when Pods are unable to communicate with each other, a service, or the outside world. Common network issues include DNS resolution problems, service misconfiguration, or blocked ports.

- Test Connectivity: Use `kubectl exec` to ping other Pods or services to verify connectivity.
- Check Network Policies: Ensure that Kubernetes network policies are not unintentionally blocking traffic.
- Check DNS Resolution: If Pods cannot resolve domain names, it might indicate a DNS issue. Check the CoreDNS logs using `kubectl logs -n kube-system`.

6. Conclusion

Debugging is a crucial skill for a Kubernetes Application Developer (CKAD). Kubernetes provides several powerful tools and commands like `kubectl describe`, `kubectl logs`, `kubectl exec`, and `kubectl top`

to help you troubleshoot and identify issues with Pods, containers, and other resources.

By familiarizing yourself with these tools and techniques, you will be better equipped to identify, diagnose, and resolve issues in your Kubernetes-based applications, ensuring smooth operation in a cloud-native environment.

